

```
#R.CHARAN SATHVIK
#21BCE7365
# VITAP MORNING SLOT
# ASSIGNMENT-4
# Data Preprocessing on Employee Attrition DataSet.
```

Import libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
a=pd.read_csv("/content/drive/MyDrive/DATASETS/WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

a

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educa
0	41	Yes	Travel_Rarely	1102	Sales		1
1	49	No	Travel_Frequently	279	Research & Development		8
2	37	Yes	Travel_Rarely	1373	Research & Development		2
3	33	No	Travel_Frequently	1392	Research & Development		3
4	27	No	Travel_Rarely	591	Research & Development		2
...	...	...	...	...	...		...
1465	36	No	Travel_Frequently	884	Research & Development		23
1466	39	No	Travel_Rarely	613	Research & Development		6
1467	27	No	Travel_Rarely	155	Research & Development		4
1468	49	No	Travel_Frequently	1023	Sales		2
1469	34	No	Travel_Rarely	628	Research & Development		8

1470 rows × 35 columns

Read the data types

a.dtypes

Age	int64
Attrition	object
BusinessTravel	object
DailyRate	int64
Department	object
DistanceFromHome	int64
Education	int64
EducationField	object
EmployeeCount	int64
EmployeeNumber	int64
EnvironmentSatisfaction	int64
Gender	object
HourlyRate	int64
JobInvolvement	int64
JobLevel	int64
JobRole	object
JobSatisfaction	int64
MaritalStatus	object
MonthlyIncome	int64
MonthlyRate	int64
NumCompaniesWorked	int64

```

Over18          object
OverTime        object
PercentSalaryHike  int64
PerformanceRating int64
RelationshipSatisfaction int64
StandardHours    int64
StockOptionLevel int64
TotalWorkingYears int64
TrainingTimesLastYear int64
WorkLifeBalance  int64
YearsAtCompany   int64
YearsInCurrentRole int64
YearsSinceLastPromotion int64
YearsWithCurrManager int64
dtype: object

```

Shape of the dataset

a.shape

```
(1470, 35)
```

Information about the dataset

a.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Age                 1470 non-null  int64
 1   Attrition           1470 non-null  object
 2   BusinessTravel      1470 non-null  object
 3   DailyRate           1470 non-null  int64
 4   Department          1470 non-null  object
 5   DistanceFromHome    1470 non-null  int64
 6   Education            1470 non-null  int64
 7   EducationField       1470 non-null  object
 8   EmployeeCount        1470 non-null  int64
 9   EmployeeNumber       1470 non-null  int64
10   EnvironmentSatisfaction 1470 non-null  int64
11   Gender              1470 non-null  object
12   HourlyRate           1470 non-null  int64
13   JobInvolvement       1470 non-null  int64
14   JobLevel             1470 non-null  int64
15   JobRole              1470 non-null  object
16   JobSatisfaction      1470 non-null  int64
17   MaritalStatus        1470 non-null  object
18   MonthlyIncome         1470 non-null  int64
19   MonthlyRate          1470 non-null  int64
20   NumCompaniesWorked   1470 non-null  int64
21   Over18               1470 non-null  object
22   OverTime             1470 non-null  object
23   PercentSalaryHike    1470 non-null  int64
24   PerformanceRating    1470 non-null  int64
25   RelationshipSatisfaction 1470 non-null  int64
26   StandardHours        1470 non-null  int64
27   StockOptionLevel     1470 non-null  int64
28   TotalWorkingYears    1470 non-null  int64
29   TrainingTimesLastYear 1470 non-null  int64
30   WorkLifeBalance      1470 non-null  int64
31   YearsAtCompany        1470 non-null  int64
32   YearsInCurrentRole   1470 non-null  int64
33   YearsSinceLastPromotion 1470 non-null  int64
34   YearsWithCurrManager 1470 non-null  int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

Statistics about the dataset

a.describe()

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	Employee
<b>count</b>	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470
<b>mean</b>	36.923810	802.485714	9.192517	2.912925	1.0	1024
<b>std</b>	9.135373	403.509100	8.106864	1.024165	0.0	602
<b>min</b>	18.000000	102.000000	1.000000	1.000000	1.0	1

Null values identification

<b>50%</b>	36.000000	802.000000	7.000000	3.000000	1.0	1024
------------	-----------	------------	----------	----------	-----	------

a.isnull().any()

```
Age                False
Attrition          False
BusinessTravel     False
DailyRate          False
Department         False
DistanceFromHome   False
Education           False
EducationField     False
EmployeeCount      False
EmployeeNumber     False
EnvironmentSatisfaction False
Gender             False
HourlyRate         False
JobInvolvement     False
JobLevel           False
JobRole            False
JobSatisfaction    False
MaritalStatus      False
MonthlyIncome      False
MonthlyRate        False
NumCompaniesWorked False
Over18             False
OverTime           False
PercentSalaryHike  False
PerformanceRating  False
RelationshipSatisfaction False
StandardHours      False
StockOptionLevel   False
TotalWorkingYears  False
TrainingTimesLastYear False
WorkLifeBalance    False
YearsAtCompany     False
YearsInCurrentRole False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool
```

a.isnull().sum()

```
Age                0
Attrition          0
BusinessTravel     0
DailyRate          0
Department         0
DistanceFromHome   0
Education           0
EducationField     0
EmployeeCount      0
EmployeeNumber     0
EnvironmentSatisfaction 0
Gender             0
HourlyRate         0
JobInvolvement     0
JobLevel           0
JobRole            0
JobSatisfaction    0
MaritalStatus      0
MonthlyIncome      0
MonthlyRate        0
NumCompaniesWorked 0
Over18             0
OverTime           0
PercentSalaryHike  0
PerformanceRating  0
RelationshipSatisfaction 0
StandardHours      0
StockOptionLevel   0
TotalWorkingYears  0
TrainingTimesLastYear 0
WorkLifeBalance    0
YearsAtCompany     0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
dtype: int64
```

```
YearsWithCurrManager      0
dtype: int64

# there are no null values

Data Visualization

d=a.corr()
d

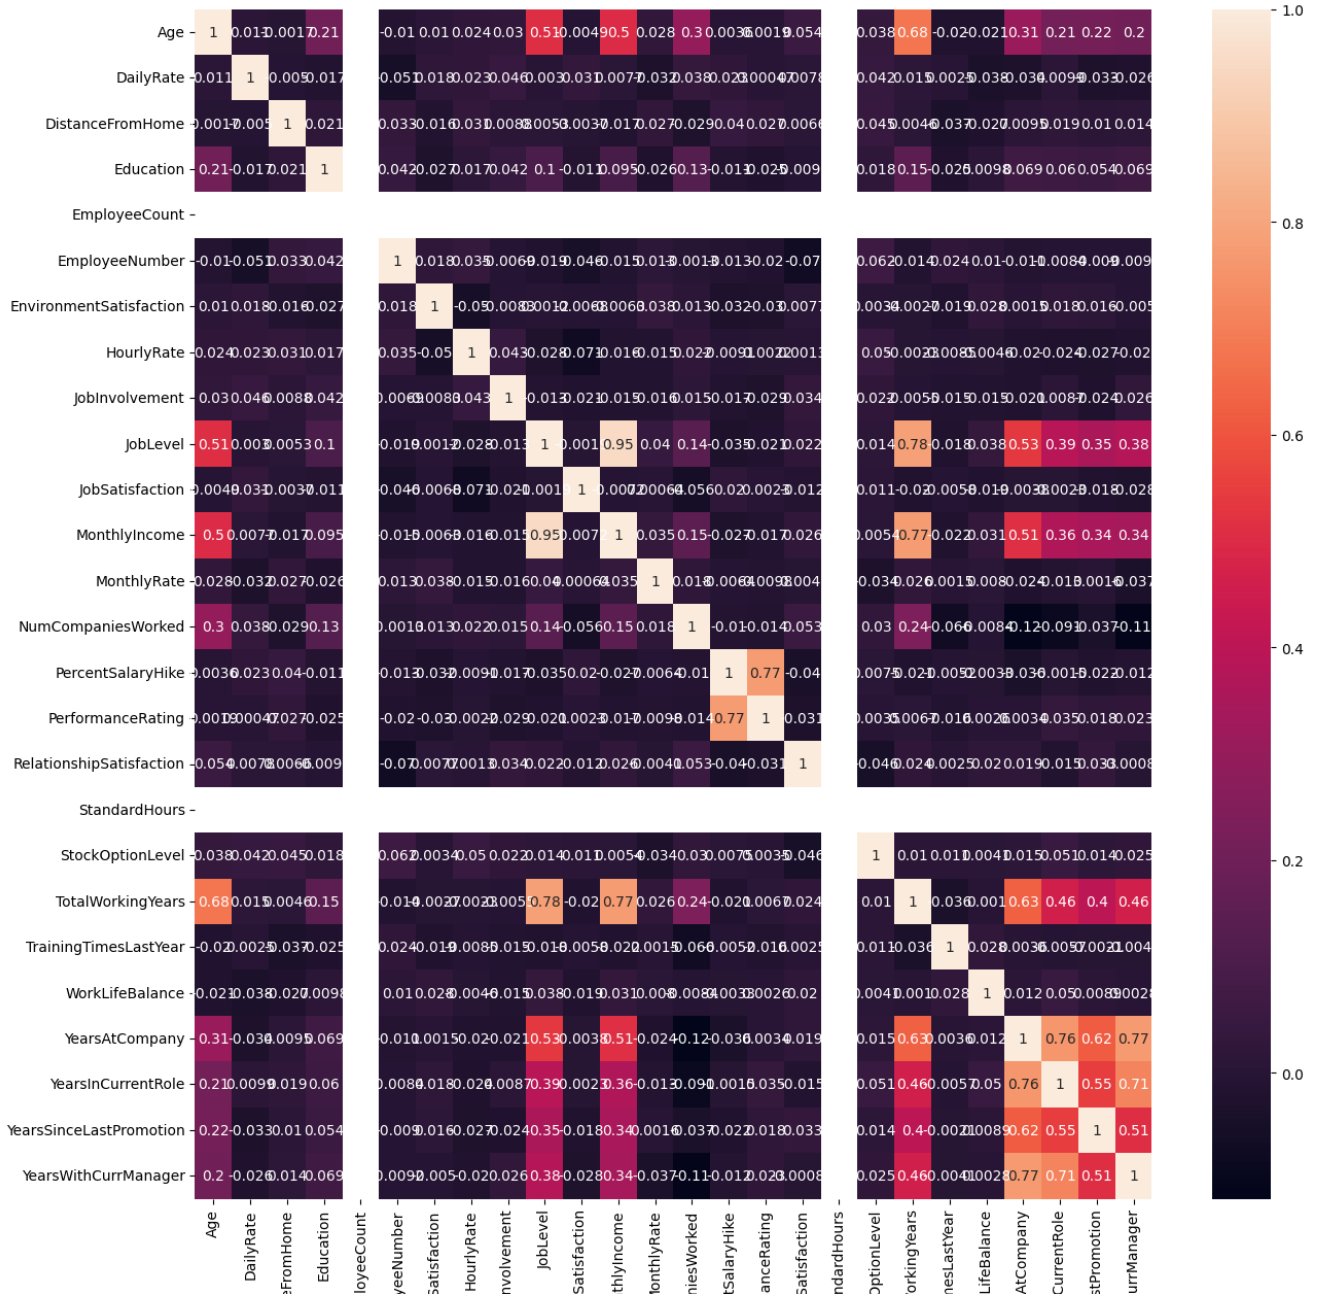
<ipython-input-12-385900cf86c7>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ve
d=a.corr()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate
Age	1.000000	0.010661	-0.001686	0.208034	NaN	-0.010145	0.010146	0.024287
DailyRate	0.010661	1.000000	-0.004985	-0.016806	NaN	-0.050990	0.018355	0.023381
DistanceFromHome	-0.001686	-0.004985	1.000000	0.021042	NaN	0.032916	-0.016075	0.031131
Education	0.208034	-0.016806	0.021042	1.000000	NaN	0.042070	-0.027128	0.016775
EmployeeCount	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.035179
EmployeeNumber	-0.010145	-0.050990	0.032916	0.042070	NaN	1.000000	0.017621	-0.049857
EnvironmentSatisfaction	0.010146	0.018355	-0.016075	-0.027128	NaN	0.017621	1.000000	-0.006888
HourlyRate	0.024287	0.023381	0.031131	0.016775	NaN	0.035179	-0.049857	1.000000
JobInvolvement	0.029820	0.046135	0.008783	0.042438	NaN	-0.006888	-0.008278	0.001212
JobLevel	0.509604	0.002966	0.005303	0.101589	NaN	-0.018519	0.001212	-0.006784
JobSatisfaction	-0.004892	0.030571	-0.003669	-0.011296	NaN	-0.046247	-0.006784	-0.006259
MonthlyIncome	0.497855	0.007707	-0.017014	0.094961	NaN	-0.014829	-0.006259	0.037600
MonthlyRate	0.028051	-0.032182	0.027473	-0.026084	NaN	0.012648	0.037600	-0.029548
NumCompaniesWorked	0.299635	0.038153	-0.029251	0.126317	NaN	-0.001251	0.012594	-0.031701
PercentSalaryHike	0.003634	0.022704	0.040235	-0.011111	NaN	-0.012944	-0.031701	-0.020359
PerformanceRating	0.001904	0.000473	0.027110	-0.024539	NaN	-0.020359	-0.029548	0.007665
RelationshipSatisfaction	0.053535	0.007846	0.006557	-0.009118	NaN	-0.069861	0.007665	NaN
StandardHours	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.037510
StockOptionLevel	0.037510	0.042143	0.044872	0.018422	NaN	0.062227	0.003432	0.042143
TotalWorkingYears	0.680381	0.014515	0.004628	0.148280	NaN	-0.014365	-0.002693	0.014515
TrainingTimesLastYear	-0.019621	0.002453	-0.036942	-0.025100	NaN	0.023603	-0.019359	-0.037848
WorkLifeBalance	-0.021490	-0.037848	-0.026556	0.009819	NaN	0.010309	0.027627	-0.034055
YearsAtCompany	0.311309	-0.034055	0.009508	0.069114	NaN	-0.011240	0.001458	0.009932
YearsInCurrentRole	0.212901	0.009932	0.018845	0.060236	NaN	-0.008416	0.018007	0.016513
YearsSinceLastPromotion	0.216513	-0.033229	0.010029	0.054254	NaN	-0.009019	0.016194	-0.026363
YearsWithCurrManager	0.202089	-0.026363	0.014406	0.069065	NaN	-0.009197	-0.004999	

26 rows x 26 columns

```
plt.subplots(figsize=(15,15))
sns.heatmap(d,annot=True)
```

&lt;Axes: &gt;



```
f = plt.figure()
f.set_figwidth(15)
f.set_figheight(12)

# Subplot 1
plt.subplot(3, 3, 1)
sns.countplot(x="Attrition", data=a)

# Subplot 2
plt.subplot(3, 3, 2)
sns.countplot(x="BusinessTravel", data=a)

# Subplot 5
plt.subplot(3, 3, 3)
sns.countplot(x="Department", data=a)


# Subplot 8
plt.subplot(3, 3, 4)
sns.countplot(x="EducationField", data=a)

# Subplot 9
plt.subplot(3, 3, 5)
sns.countplot(x="OverTime", data=a)

# Adjust layout
```

```
plt.tight_layout()


# Show the plots
plt.show()
```



Outlier Detection


```
sns.boxplot(x="Age",data=a)

<Axes: xlabel='Age'>
```




```
sns.boxplot(x="DailyRate",data=a)
```

<Axes: xlabel='DailyRate'>




```
sns.boxplot(x="DistanceFromHome",data=a)
```

<Axes: xlabel='DistanceFromHome'>




```
sns.boxplot(x="Education",data=a)
```

<Axes: xlabel='Education'>




```
sns.boxplot(x="EmployeeCount",data=a)
```

```
<Axes: xlabel='EmployeeCount'>
```




```
sns.boxplot(x="EnvironmentSatisfaction",data=a)
```

```
<Axes: xlabel='EnvironmentSatisfaction'>
```



```
sns.boxplot(x="HourlyRate",data=a)
```

```
<Axes: xlabel='HourlyRate'>
```



```
# there are no outliers , the data is clean
```

```
Splitting dependent and independent variables
```



```
x=a.drop(columns=["Attrition"],axis=1)
x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environment
0	41	Travel_Rarely	1102	Sales		1	Life Sciences	1		1
1	49	Travel_Frequently	279	Research & Development		8	Life Sciences	1		2
2	37	Travel_Rarely	1373	Research & Development		2	Other	1		4
3	33	Travel_Frequently	1392	Research & Development		3	Life Sciences	1		5
4	27	Travel_Rarely	591	Research & Development		2	Medical	1		7

5 rows × 34 columns

```
x.shape
```

(1470, 34)

```
y=a["Attrition"]
y.head()
```

```
0    Yes
1     No
2     Yes
3     No
4     No
Name: Attrition, dtype: object
```

```
y.shape
```

(1470,)

Encoding

```
from sklearn.preprocessing import LabelEncoder
```

```
l=LabelEncoder()
```

```
x["Gender"]=l.fit_transform(x["Gender"])
x['Gender']
```

```
0      0
1      1
2      1
3      0
4      1
..
1465   1
1466   1
1467   1
1468   1
1469   1
Name: Gender, Length: 1470, dtype: int64
```

```
x['Gender'].value_counts()
```

```
1    882
0    588
Name: Gender, dtype: int64
```

```
x['Gender'].nunique()
```

2

```
x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environment
0	41	Travel_Rarely	1102	Sales		1	2	Life Sciences	1	1
1	49	Travel_Frequently	279	Research & Development		8	1	Life Sciences	1	2
2	37	Travel_Rarely	1373	Research & Development		2	2	Other	1	4
3	33	Travel_Frequently	1392	Research & Development		3	4	Life Sciences	1	5

Research &

```
Dept = pd.get_dummies(a, columns=["Department"])
print(Dept)
```

1466	4	...	9	5
1467	2	...	6	0
1468	4	...	17	3
1469	2	...	6	3

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
0	1	6	4	
1	3	10	7	
2	3	0	0	
3	3	8	7	
4	3	2	2	
...	...	...	...	
1465	3	5	2	
1466	3	7	7	
1467	3	6	2	
1468	2	9	6	
1469	4	4	3	

	YearsSinceLastPromotion	YearsWithCurrManager	\
0	0	5	
1	1	7	
2	0	0	
3	3	0	
4	2	2	
...	...	...	
1465	0	3	
1466	1	7	
1467	0	3	
1468	0	8	
1469	1	2	

	Department_Human Resources	Department_Research & Development	\
0	0	0	
1	0	1	
2	0	1	
3	0	1	
4	0	1	
...	...	...	
1465	0	1	
1466	0	1	
1467	0	1	
1468	0	0	
1469	0	1	

	Department_Sales
0	1
1	0
2	0
3	0
4	0
...	...
1465	0
1466	0
1467	0
1468	1
1469	0

[1470 rows x 37 columns]

```
print(x)
```

	Age	BusinessTravel	DailyRate	Department	\
0	41	Travel_Rarely	1102	Sales	
1	49	Travel_Frequently	279	Research & Development	
2	37	Travel_Rarely	1373	Research & Development	
3	33	Travel_Frequently	1392	Research & Development	
4	27	Travel_Rarely	591	Research & Development	
...	...	...	...	...	
1465	36	Travel_Frequently	884	Research & Development	
1466	39	Travel_Rarely	613	Research & Development	
1467	27	Travel_Rarely	155	Research & Development	
1468	49	Travel_Frequently	1023	Sales	

1469	34	Travel_Rarely	628	Research & Development	
	DistanceFromHome	Education	EducationField	EmployeeCount	\
0	1	2	Life Sciences	1	
1	8	1	Life Sciences	1	
2	2	2	Other	1	
3	3	4	Life Sciences	1	
4	2	1	Medical	1	
...	...	...	...	...	
1465	23	2	Medical	1	
1466	6	1	Medical	1	
1467	4	3	Life Sciences	1	
1468	2	3	Medical	1	
1469	8	3	Medical	1	
	EmployeeNumber	EnvironmentSatisfaction	...	RelationshipSatisfaction	\
0	1	2	...	1	
1	2	3	...	4	
2	4	4	...	2	
3	5	4	...	3	
4	7	1	...	4	
...	...	...	...	...	
1465	2061	3	...	3	
1466	2062	4	...	1	
1467	2064	2	...	2	
1468	2065	4	...	4	
1469	2068	2	...	1	
	StandardHours	StockOptionLevel	TotalWorkingYears		\
0	80	0	8		
1	80	1	10		
2	80	0	7		
3	80	0	8		
4	80	1	6		
...	...	...	...		
1465	80	1	17		
1466	80	1	9		
1467	80	1	6		
1468	80	0	17		
1469	80	0	6		
	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany		\
0	0	1	6		
1	3	3	10		
2	3	3	0		
3	3	3	8		
4	2	2	7		

```
a.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSa
0	41	Yes	Travel_Rarely	1102	Sales	1	2	1	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	1	2	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	1	4	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	1	5	
4	27	No	Travel_Rarely	591	Research & Development	2	1	1	7	

5 rows × 40 columns

```
x.head()
```

```
Age BusinessTravel DailyRate Department DistanceFromHome Education EducationField EmployeeCount EmployeeNumber Environment
Dept=pd.get_dummies(x["Department"],drop_first=True)
Dept
```

	Research & Development	Sales
0	0	1
1	1	0
2	1	0
3	1	0
4	1	0
...	...	...
1465	1	0
1466	1	0
1467	1	0
1468	0	1
1469	1	0

1470 rows × 2 columns

```
x=pd.concat([x,Dept],axis=1)
x.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Environment
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	

5 rows × 36 columns

Feature Scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X = a[['Age', 'MonthlyIncome', 'YearsAtCompany', 'JobSatisfaction', 'EnvironmentSatisfaction', 'YearsWithCurrManager', 'WorkLifeBalance']]
Y = a['Attrition']

X.head()
```

	Age	MonthlyIncome	YearsAtCompany	JobSatisfaction	EnvironmentSatisfaction	YearsWithCurrManager	WorkLifeBalance
0	41	5993	6	4	2	5	1
1	49	5130	10	2	3	7	3
2	37	2090	0	3	4	0	3
3	33	2909	8	3	4	0	3
4	27	3468	2	2	1	2	3

```
x.tail()
```

x

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	Envir
1465	36	Travel_Frequently	884	Research & Development	23	2	Medical	1	2061	
1466	39	Travel_Rarely	613	Research & Development	6	1	Medical	1	2062	
1467	27	Travel_Rarely	155	Research & Development	4	3	Life Sciences	1	2064	
1468	49	Travel_Frequently	1023	Sales	2	3	Medical	1	2065	
1469	34	Travel_Rarely	628	Research & Development	8	3	Medical	1	2068	
...										
0	41	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	Other	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	Medical	1	7	
...	...	...	...	...	...	...	...	...	...	...
1465	36	Travel_Frequently	884	Research & Development	23	2	Medical	1	2061	
1466	39	Travel_Rarely	613	Research & Development	6	1	Medical	1	2062	
1467	27	Travel_Rarely	155	Research & Development	4	3	Life Sciences	1	2064	
1468	49	Travel_Frequently	1023	Sales	2	3	Medical	1	2065	
1469	34	Travel_Rarely	628	Research & Development	8	3	Medical	1	2068	
1470 rows x 36 columns										

Splitting data into test and train

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

X_train,X_test,Y_train,Y_test.shape
```

(	Age	MonthlyIncome	YearsAtCompany	JobSatisfaction	\
1097	24	2296	1	1	
727	18	1051	0	4	
254	29	6931	3	4	
1175	39	5295	5	2	
1341	31	4197	10	3	
...	...	...	...	...	
1130	35	3407	10	3	
1294	41	6870	3	2	
860	22	2853	0	4	
1459	29	4025	4	2	
1126	50	19331	1	3	
	EnvironmentSatisfaction	YearsWithCurrManager	WorkLifeBalance		
1097	3	0	3		
727	2	0	3		
254	4	2	3		
1175	4	0	3		
1341	2	2	3		
...	...	...	...		
1130	2	8	2		
1294	2	2	1		
860	3	0	3		
1459	4	3	3		
1126	3	0	3		
[1176 rows x 7 columns],					

	EnvironmentSatisfaction	YearsWithCurrManager	WorkLifeBalance
1041	4	3	3
184	4	3	3
1222	4	0	3
67	2	0	3
220	4	7	4
...	...	...	...
567	4	4	3
560	2	0	3
945	4	2	3
522	4	0	3
651	3	7	3

## Logistic Regression

## Model Building & Import the model building Libraries

```
model.fit(X_train, Y_train)
```

```
pred=model.predict(X_test)
```

pred

[illegible]

1041	No
184	No
1222	Yes
67	No
220	No

```
...
567     No
568     No
945     No
522     No
651     No
Name: Attrition, Length: 294, dtype: object

a

Age Attrition BusinessTravel DailyRate Department DistanceFromHome Education EmployeeCount EmployeeNumber Environmen
0 41 Yes Travel_Rarely 1102 Sales 1 2 1 1
1 49 No Travel_Frequently 279 Research & Development 8 1 1 2
2 37 Yes Travel_Rarely 1373 Research & Development 2 2 1 4
3 33 No Travel_Frequently 1392 Research & Development 3 4 1 5
4 27 No Travel_Rarely 591 Research & Development 2 1 1 7
... ..
1465 36 No Travel_Frequently 884 Research & Development 23 2 1 2061
1466 39 No Travel_Rarely 613 Research & Development 6 1 1 2062
1467 27 No Travel_Rarely 155 Research & Development 4 3 1 2064
1468 49 No Travel_Frequently 1023 Sales 2 3 1 2065
1469 34 No Travel_Rarely 628 Research & Development 8 3 1 2068

1470 rows x 40 columns
```

Evaluation of classification model

```
#Accuracy score
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,roc_auc_score,roc_curve

accuracy = accuracy_score(Y_test, pred)

report = classification_report(Y_test, pred, zero_division=1)

print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')

Accuracy: 0.8673469387755102
Classification Report:
              precision    recall  f1-score   support

     No       0.87        1.00        0.93        255
     Yes       1.00        0.00        0.00         39

 accuracy         0.87        0.87        0.87        294
 macro avg       0.93        0.50        0.46        294
 weighted avg     0.88        0.87        0.81        294

confusion_matrix(Y_test,pred)

array([[255,  0],
       [ 39,  0]])

pd.crosstab(Y_test,pred)
```

col\_0 No 

Roc-AUC curve

fpr tpr

probability=model.predict\_proba(X\_test)[: ,1]

probability

```

0.14963007, 0.15969356, 0.20644099, 0.08193936, 0.18537088,
0.16096129, 0.02189805, 0.15660552, 0.11782876, 0.18248771,
0.13287268, 0.14334387, 0.0892007 , 0.06858367, 0.05708061,
0.1753651 , 0.14395111, 0.10012064, 0.15057687, 0.2329628 ,
0.03338823, 0.27116899, 0.15771847, 0.18762417, 0.10029771,
0.10548668, 0.15048832, 0.12644386, 0.14778903, 0.2030313 ,
0.06737083, 0.04935137, 0.35253675, 0.19926437, 0.23846212,
0.08198467, 0.28864726, 0.23955634, 0.19282515, 0.22246873,
0.11288909, 0.17545014, 0.24051176, 0.14059822, 0.32377579,
0.08977525, 0.15148043, 0.01896052, 0.14635136, 0.20158982,
0.10191406, 0.10573264, 0.08537077, 0.1631479 , 0.12443613,
0.10510977, 0.33623452, 0.11027653, 0.05493965, 0.28005007,
0.18450873, 0.12499531, 0.17197795, 0.17873294, 0.06110176,
0.18127058, 0.08791989, 0.15005295, 0.15959692, 0.19866202,
0.07388538, 0.19341696, 0.19100387, 0.08712656, 0.08033949,
0.02928375, 0.13253218, 0.05956382, 0.16844953, 0.08753921,
0.17957672, 0.12899389, 0.16872069, 0.16947305, 0.12397644,
0.1099147 , 0.24576674, 0.07821105, 0.2716565 , 0.12140547,
0.06524951, 0.1337184 , 0.14536957, 0.18726004, 0.10915274,
0.04570312, 0.10169758, 0.07390408, 0.22704117, 0.07208355,
0.08035364, 0.18593691, 0.16647288, 0.10818369, 0.05315879,
0.17696614, 0.18973955, 0.22476227, 0.17342537, 0.21403334,
0.16943373, 0.16771766, 0.09747364, 0.11387728, 0.2559594 ,
0.32393512, 0.08431327, 0.13118746, 0.10751731, 0.09837008,
0.25991497, 0.18954525, 0.11954205, 0.10534474, 0.09694665,
0.07268098, 0.30507638, 0.06501248, 0.14080365, 0.1255734 ,
0.11537899, 0.23299235, 0.17264787, 0.24765337, 0.06927027,
0.21512755, 0.09901074, 0.16646941, 0.08047622, 0.03233445,
0.15363939, 0.14131117, 0.25851265, 0.26761484, 0.1665985 ,
0.10685997, 0.11549038, 0.19827264, 0.19076354, 0.13247131,
0.26173972, 0.17180386, 0.21324175, 0.04115976, 0.15054569,
0.16012435, 0.09434315, 0.09921354, 0.22000675, 0.06421677,
0.16643204, 0.12016002, 0.14827189, 0.08450615, 0.05725373,
0.12102272, 0.02681568, 0.18300015, 0.21076054, 0.11715199,
0.16127828, 0.18483891, 0.09043029, 0.14086669, 0.20253644,
0.0594472 , 0.10383826, 0.01617733, 0.15428555, 0.08595314,
0.22434066, 0.11577713, 0.07998958, 0.07811109, 0.12006351,
0.12845942, 0.14824842, 0.10405812, 0.19816497, 0.1162661 ,
0.21477996, 0.24395257, 0.04972863, 0.2156586 , 0.16831872,
0.17867722, 0.15398516, 0.21871738, 0.03416769, 0.07072713,
0.22242289, 0.10244091, 0.10919764, 0.12517809, 0.0706504 ,
0.07399615, 0.24438034, 0.17159597, 0.17617076, 0.10663942,
0.13898632, 0.15178097, 0.10545546, 0.2723432 , 0.07462743,
0.23465253, 0.26405405, 0.10124306, 0.3028089 , 0.12410107,
0.1909214 , 0.20302625, 0.13276688, 0.0401135 , 0.18943046,
0.23129363, 0.25951761, 0.08630086, 0.21347439, 0.20469075,
0.13330949, 0.08581729, 0.10996842, 0.06690194, 0.04616928,
0.18853288, 0.11542819, 0.21231547, 0.03597583, 0.07176025,
0.17130681, 0.11593175, 0.23407496, 0.1533375 , 0.09696206,
0.16256038, 0.06366454, 0.04689748, 0.0855508 , 0.23703024,
0.07106702, 0.18067446, 0.2069784 , 0.22648723, 0.02715875,
0.17170263, 0.14167865, 0.276632 , 0.10463943, 0.12037205,
0.21133882, 0.02933273, 0.0973697 , 0.23466029, 0.23184945,
0.1882965 , 0.04906958, 0.19036583, 0.1399965 , 0.11412922,
0.22223015, 0.12517666, 0.24824295, 0.07113102, 0.07508479,
0.14609486, 0.15491467, 0.18318556, 0.09382192, 0.04811606,
0.20893659, 0.20088061, 0.23217748, 0.10747859, 0.11268901,
0.25784861, 0.07464244, 0.1744561 , 0.09272658])

```

```

from sklearn.preprocessing import LabelBinarizer
lb = LabelBinarizer()
Y_test_bin = lb.fit_transform(Y_test)
fpr, tpr, thresholds = roc_curve(Y_test_bin, probability)


```

```

plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()

```





```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
dt_model = DecisionTreeClassifier(random_state=50)
```

```
dt_model.fit(X_train, Y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=50)
```

```
dt_predictions = dt_model.predict(X_test)
```

```
dt_accuracy = accuracy_score(Y_test, dt_predictions)
```

```
dt_report = classification_report(Y_test, dt_predictions)
```

```
print(f'Decision Tree Accuracy: {dt_accuracy}')
```

```
Decision Tree Accuracy: 0.7789115646258503
```

```
print(f'Decision Tree Classification Report:\n{dt_report}')
```

```
Decision Tree Classification Report:
              precision    recall  f1-score   support

    No         0.90         0.84         0.87         255
    Yes         0.28         0.41         0.33          39

 accuracy         0.78         0.78         0.78         294
 macro avg        0.59         0.62         0.60         294
 weighted avg     0.82         0.78         0.80         294
```

## Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_model = RandomForestClassifier(random_state=50)
```

```
rf_model.fit(X_train, Y_train)
```

```
RandomForestClassifier
RandomForestClassifier(random_state=50)
```

```
rf_predictions = rf_model.predict(X_test)
```

```
rf_accuracy = accuracy_score(Y_test, rf_predictions)
```

```
rf_report = classification_report(Y_test, rf_predictions)
```

```
print(f'Random Forest Accuracy: {rf_accuracy}')
```

Random Forest Accuracy: 0.8435374149659864

```
print(f'Random Forest Classification Report:\n{rf_report}')
```

Random Forest Classification Report:

	precision	recall	f1-score	support
No	0.88	0.95	0.91	255
Yes	0.33	0.18	0.23	39
accuracy			0.84	294
macro avg	0.61	0.56	0.57	294
weighted avg	0.81	0.84	0.82	294